The Economist's Guide to [Prompt Engineering]

Extracting `Alpha` from Language Models

NotebookLM

The Paradigm Shift: Programming in Natural Language

The Old Syntax

```
df.groupby('region')
['sales'].mean()
.sort_values(ascendi
ng=False)
```

Rigid Syntax.
High barrier to entry.

The New Syntax

Calculate the average sales per region and sort them in descending order.

Natural Language.
Boundless utility.

The Output

Region	Average Sales
North America	50,000 USD
Europe	45,000 USD
Asia-Pacific	40,000 USD

Large Language Models (LLMs) are **text prediction engines**, not databases. They do not have built-in task lists. Your prompt is the code that compiles their next action and bounds their output space.

NotebookLM

The Cost of Ambiguity: Vague vs. Steered Inputs

The Vague Prompt	The Steered Prompt
Input: `Write an ad about our product.`	Input: `You are a senior B2B copywriter. Write a 2-sentence LinkedIn ad for our project management SaaS, a direct competitor to Asana. The audience is ops managers at midsize companies. The tone is confident but not salesy. End with a clear CTA.`
Output Quality: Generic marketing fluff, wrong length, wrong tone, zero actionability. The model is forced to guess your intent.	Output Quality: On-brand, specifically scoped, highly actionable.
Resource Waste: High (requires manual rewriting).	Resource Waste: Zero.

Calculated ROI: In prompt engineering, ambiguity is an expensive tax. Precision yields high-utility assets.

NotebookLM

The Core Algorithm: T-C-R-E-I

T

Task: What exactly do you want the AI to do? (The core objective).

C

Context: What background information does the model need to succeed? (The variables).

R

References: What examples or baseline data can you provide? (The training data).

E

Evaluate: Does the output match your precise expectations? (The diagnostic).

I

Iterate: Refine the prompt to close the gap between output and expectation. (The loop).

Remember the sequence: **T**iny **C**rabs **R**ide **E**normous **I**guanas.

NotebookLM

Bounding the Output Space: The Big Four Constraints

```
You are a Chief Economist.

Summarize this earnings report
for retail investors. Keep the
language accessible but
analytically rigorous, avoiding
jargon. Output the insights as a
three-column markdown table:
Metric, Previous Quarter,
Projected Quarter.
```

Role: Defines the "who." Sets the baseline expertise and knowledge parameters of the model.

Audience: Defines the "to whom." Calibrates the complexity and vocabulary of the output.

Tone: Defines the "how." Establishes the emotional and professional register.

Format: Defines the "what." Forces structured data (JSON, Markdown, bullet lists) instead of generic prose.

NotebookLM

The Economics of Iteration (Closing the 20% Gap)

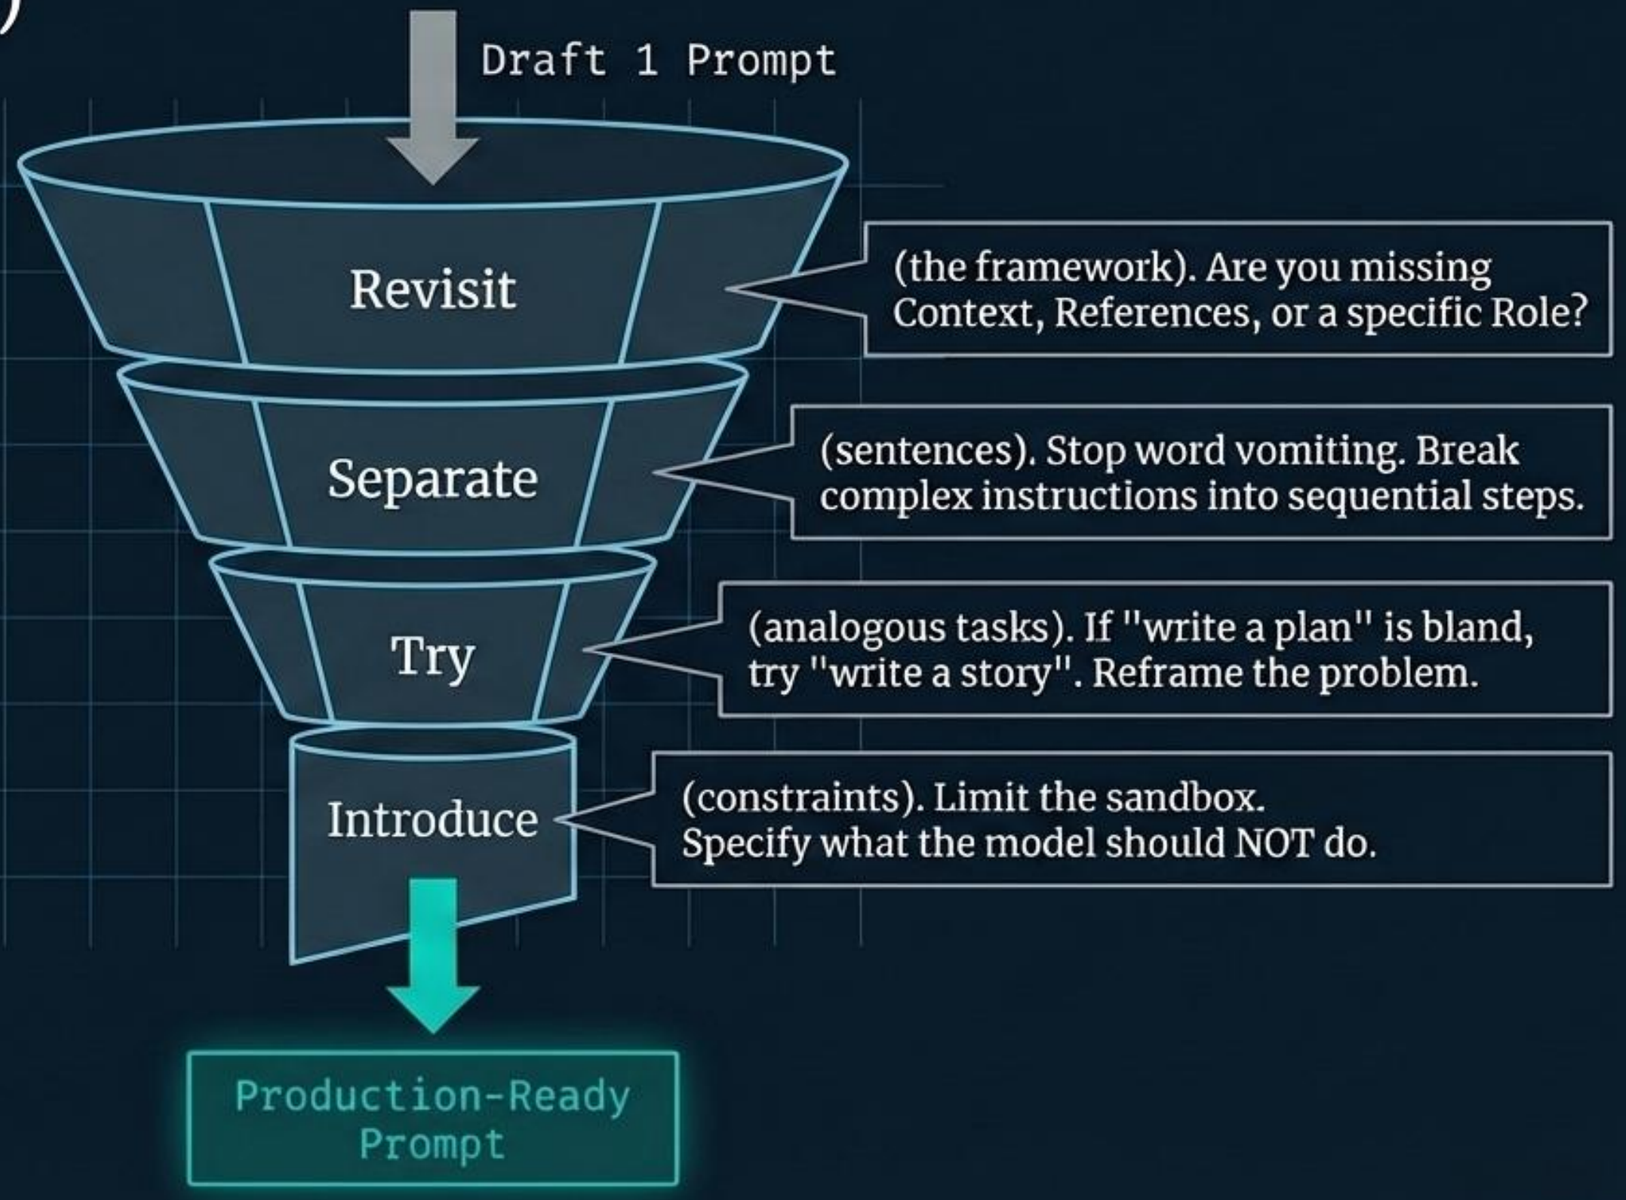


```
Mnemonic Anchor:
Refine Structure
Test
Iterate.
```

NotebookLM

Applied Economics: Grocery Store Data Analysis

The Direct Command

Prompt:

Attached is a Google sheet of store data. Create a new column that calculates the average sales per customer for each store.

Result:

Generates the exact spreadsheet formula required, eliminating manual syntax errors.

The Analytical Steer

Prompt:

Act as a retail analyst. Give me insights into the relationship between daily customer count, items available, and total sales based on the given data.

Result:

Identifies non-obvious trends (e.g., revealing no clear correlation between inventory size and total sales, prompting further strategic investigation).

NotebookLM

Advanced Toolkit I: Few-Shot Prompting

Concept: Providing the model with a few examples of desired input-output pairs before giving it the actual task trains it to recognize patterns and lock into strict formatting constraints.

Pattern Recognition

Training Example 1
Input: "Customer complaint about high prices"
→ Output: [SENTIMENT_NEG] Price_complaint

Training Example 2
Input: "Request for quarterly sales report"
→ Output: [DARA_REQ] Quarterly_report

Training Example 3
Input: "Survey : Store layout is confusing."
-> Output: [FEEDBACK_NEG] Store_layout

Target Task
Input: "Analyst note: Competitors lowered prices of milk"

Neural Processor

Guaranteed Output:
[MARKET_INTEL] Comptetitor_pricing

Takeaway: Few-shot prompting is essential when classifying qualitative data (like economic sentiment, support tickets, or survey responses) into rigid, database-ready categories.

NotebookLM

Advanced Toolkit II: Chain of Thought (CoT)

The Syntax Trigger: Append the phrase: "Explain your thought process step-by-step before answering."

The Problem: LLMs are text predictors, not calculators. Asked to solve multi-step financial logic immediately (Zero-Shot), they frequently guess the wrong answer.

Correct Output

```
[RESULT: $4,250,000.00
- Validated]
```

Complex Financial Problem

```
Calculate the net present value
(NPV) of a 5-year project with
varying cash flows, a declining
discount rate, and terminal
value assumptions.
...
```

The Solution: By generating intermediate steps, the model feeds its own correct logic back into its context window, drastically reducing errors in complex economic scenarios.

NotebookLM

System Guardrails & Analytical Risks

Risk 1: Hallucinations (Confidence ≠ Accuracy)

The Threat

When a generative AI tool provides outputs that are **inconsistent, mathematically incorrect,** or entirely nonsensical, but presents them with **absolute certainty.**

The Mitigation

Human in the Loop. Never treat an LLM as an **oracle.** Treat it as a highly capable, slightly **reckless junior analyst.** Always verify outputs against primary data.

[ACTION: MANUAL_AUDIT]

Risk 2: Data Privacy & Security

The Threat

Public models train on user inputs. Pasting proprietary data into a public prompt exposes it permanently.

The Mitigation

Zero-Trust Input. Never input proprietary financial models, Personally Identifiable Information (PII), or unreleased corporate data into public consumer models.

[PROTOCOL: DATA_SANITIZATION]

NotebookLM

The Business Prompting Cheat Sheet

[1] The Core Algorithm (TCREI)

- **T | Task:** Define the core objective.
- **C | Context:** Provide the variables and background.
- **R | References:** Supply baseline data.
- **E | Evaluate:** Diagnose the output.
- **I | Iterate:** Close the gap.

[2] Bounding the Output (The Big 4)

Role: Who is the AI?

Audience: Who is receiving this?

Tone: What is the register?

Format: How is the data structured?

[3] Iteration Protocols (RSTI)

- **R | Revisit:** Check the framework.
- **S | Separate:** Break up complex sentences.
- **T | Try Analogous:** Reframe the task entirely.
- **I | Introduce Constraints:** Tell it what *not* to do.

[4] Advanced Triggers

Few-Shot: Provide Input/Output examples for perfect data classification.

Chain of Thought: Append "**Explain your reasoning step-by-step**" to force mathematical/logical accuracy.

NotebookLM